

NiFi-Example-Workflow

[Introduction](#)

[Prerequisites](#)

[Step 1 - Cloning the repository](#)

[Step 2 - Setting up kafka server](#)

[Step 3 - Produce data to Kafka](#)

[Step 4 - Consume kafka to NiFi](#)

[Step 5 - Extract the contents of flowfile into attributes](#)

[Step 6 - Update the attributes of flowfile to add processedAtTimestamp](#)

[Step 7 - Rename and pull the values back to flowfile contents using Jolt processor](#)

[Step 8 - Route the flowfiles to different destination based on attributes](#)

[Step 9 - Configure MySQL and PostgreSQL database connection](#)

Introduction

This document is for sending data to MySQL from NiFi. This workflow is designed for covering multiple processors.

It talks about consuming data from kafka and send it to multiple destinations like MySQL, PostgreSQL, ElasticSearch, Kafka etc.

Prerequisites

Make sure python3, pip3, git, docker and docker compose are installed on your system.

If not, run as root the following on Ubuntu based system

```
# Add Docker's official GPG key:
sudo apt-get update
sudo apt-get install ca-certificates curl
sudo install -m 0755 -d /etc/apt/keyrings
sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg -o
sudo chmod a+r /etc/apt/keyrings/docker.asc
```

```
# Add the repository to Apt sources:
echo \
    "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keys/
    $(. /etc/os-release && echo "$VERSION_CODENAME") stable" | \
    sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
sudo apt-get update

sudo apt-get install git python3-pip docker-ce docker-ce-cli containerd.io
```

For debian systems, follow

```
# Add Docker's official GPG key:
sudo apt-get update
sudo apt-get install ca-certificates curl
sudo install -m 0755 -d /etc/apt/keyrings
sudo curl -fsSL https://download.docker.com/linux/debian/gpg -o /etc/apt/keyrings/docker.asc
sudo chmod a+r /etc/apt/keyrings/docker.asc

# Add the repository to Apt sources:
echo \
    "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keys/
    $(. /etc/os-release && echo "$VERSION_CODENAME") stable" | \
    sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
sudo apt-get update

sudo apt-get install git python3-pip docker-ce docker-ce-cli containerd.io
```

Refer to here <https://docs.docker.com/engine/install/> for installing docker components on other systems.

Step 1 - Cloning the repository

Clone the git repository which contains everything related to this lab example.

```
git clone https://github.com/cloudpoet-in/nifi-training.git
```

Step 2 - Setting up kafka server

Setup kafka server if you dont have it already.

In git repo, navigate to directory `nifi-training/sandbox-deployments/kafka` and then run

```
docker compose up -d
```

This will start a kafka server which will listen on `localhost:9092`

Step 3 - Produce data to Kafka

Produce some data to kafka using simulator script.

Simulator script requires a python library called as `confluent-kafka` which can be installed by running

```
pip3 install confluent-kafka

# if this doesnt work, run
# pip3 install confluent-kafka --break-system-packages
```

Once dependency is ready, navigate to `nifi-training/kafka-producer-simulator` and run

```
python3 kafka_send.py
```

This will continuously produce some dummy records to kafka running at `localhost:9092` in topic named `nifi-dummy`. Keep the script running in the background as it will produce data every 5 seconds. Feel free to update it as per needs.

Successful logs looks like this

```

root@nifi-2:~/nifi-training/simulator# python3 kafka_send.py
Message produced: {"source_type": "warehouse", "source_name": "Warehouse 1", "temperature": -2.21, "pressure": 980.54, "timestamp": 1727170984}
Message produced: {"source_type": "city", "source_name": "City Center", "temperature": 3.26, "pressure": 1035.9, "timestamp": 1727170990}
Message produced: {"source_type": "warehouse", "source_name": "Warehouse 1", "temperature": -11.15, "pressure": 1009.63, "timestamp": 1727170995}
Message produced: {"source_type": "city", "source_name": "City Center", "temperature": 16.29, "pressure": 975.28, "timestamp": 1727171000}
Message produced: {"source_type": "city", "source_name": "City Center", "temperature": 34.84, "pressure": 1020.96, "timestamp": 1727171005}
Message produced: {"source_type": "factory", "source_name": "Factory A", "temperature": -15.37, "pressure": 958.01, "timestamp": 1727171010}
Message produced: {"source_type": "city", "source_name": "City Suburb", "temperature": -14.69, "pressure": 1026.51, "timestamp": 1727171015}
Message produced: {"source_type": "warehouse", "source_name": "Warehouse 2", "temperature": -9.9, "pressure": 986.02, "timestamp": 1727171020}
Message produced: {"source_type": "factory", "source_name": "Factory B", "temperature": 21.71, "pressure": 970.93, "timestamp": 1727171025}
Message produced: {"source_type": "warehouse", "source_name": "Warehouse 1", "temperature": 9.71, "pressure": 993.14, "timestamp": 1727171030}
Message produced: {"source_type": "factory", "source_name": "Factory B", "temperature": 36.37, "pressure": 953.31, "timestamp": 1727171035}
Message produced: {"source_type": "factory", "source_name": "Factory A", "temperature": 26.05, "pressure": 955.45, "timestamp": 1727171040}
Message produced: {"source_type": "city", "source_name": "City Center", "temperature": 9.52, "pressure": 967.77, "timestamp": 1727171045}
Message produced: {"source_type": "city", "source_name": "City Center", "temperature": -2.29, "pressure": 976.77, "timestamp": 1727171050}
Message produced: {"source_type": "city", "source_name": "City Suburb", "temperature": 24.37, "pressure": 1048.46, "timestamp": 1727171055}
Message produced: {"source_type": "factory", "source_name": "Factory A", "temperature": -9.72, "pressure": 978.79, "timestamp": 1727171060}

```

Step 4 - Consume kafka to NiFi

Drag a processor in nifi named **ConsumeKafka_2_6** and configure it as shown. Add unique GroupID.

Configure Processor | ConsumeKafka_2_6 1.27.0

Invalid

SETTINGS | SCHEDULING | **PROPERTIES** | RELATIONSHIPS | COMMENTS

Required field

| Property | Value |
|------------------------------|-------------------|
| Kafka Brokers | localhost:9092 |
| Topic Name(s) | nifi-dummy |
| Topic Name Format | names |
| Group ID | testing123 |
| Commit Offsets | true |
| Max Uncommitted Time | 1 secs |
| Honor Transactions | true |
| Message Demarcator | No value set |
| Separate By Key | false |
| Security Protocol | PLAINTEXT |
| SASL Mechanism | GSSAPI |
| Kerberos Credentials Service | No value set |

CANCEL APPLY

The property named **offset Reset** can be changed to consume from end of topic or start of the topic.

Configure Processor | ConsumeKafka_2_6 1.27.0

Invalid

SETTINGS

SCHEDULING

PROPERTIES

RELATIONSHIPS

COMMENTS

Required field

Property

Value

| | |
|--------------------------------------|---------------|
| Kerberos Credentials Service | No value set |
| Kerberos User Service | No value set |
| Kerberos | |
| Kerberos | |
| Kerberos | No value set |
| SSL Context | No value set |
| Key Attribute | UTF-8 Encoded |
| Offset Reset | latest |
| Message Header Encoding | UTF-8 |
| Headers to Add as Attributes (Regex) | No value set |
| Max Poll Records | 10000 |
| Communications Timeout | 60 secs |

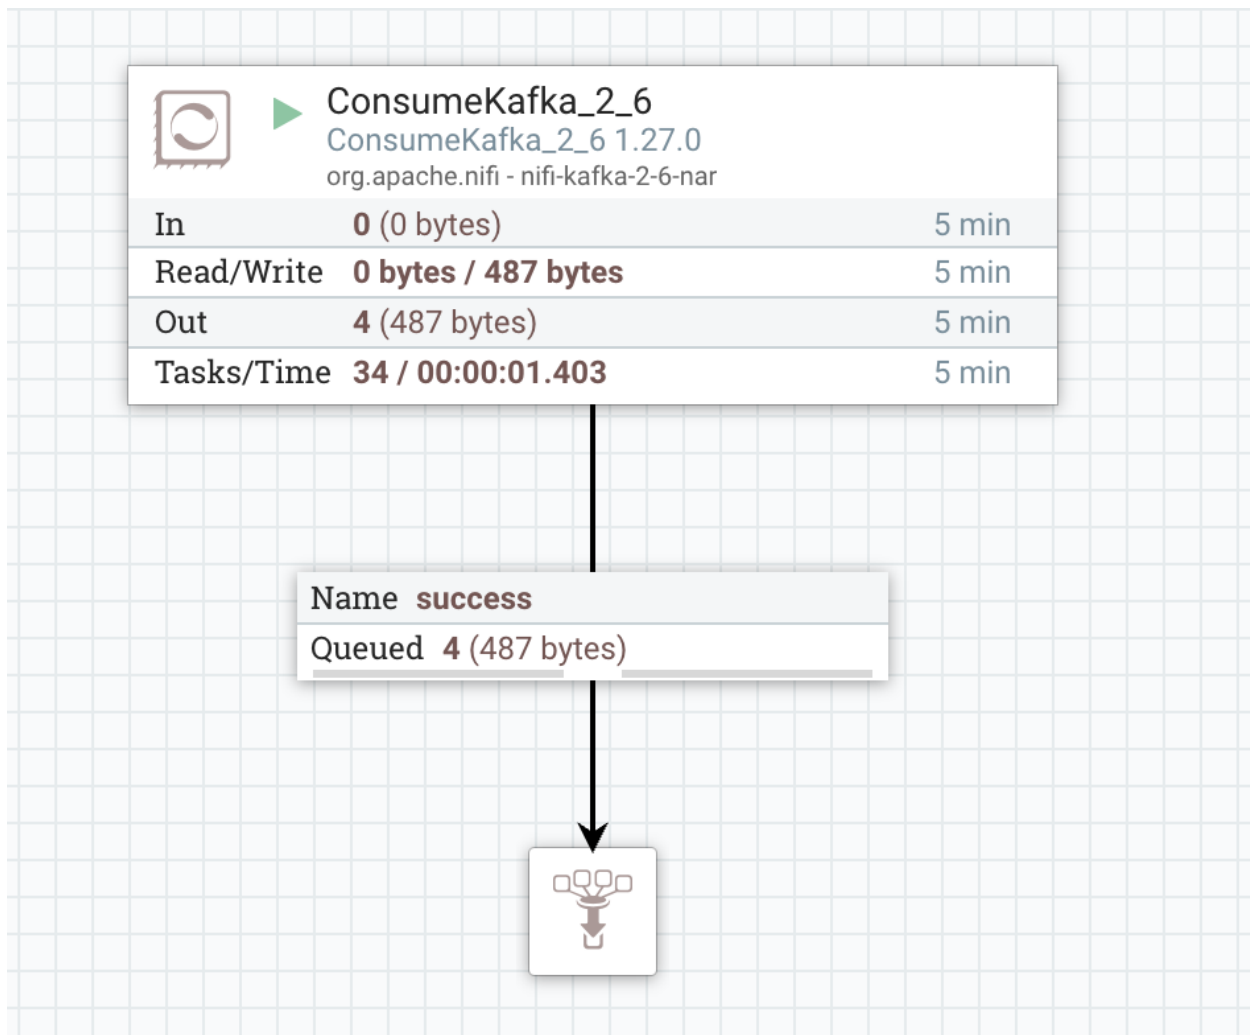
CANCEL

APPLY

Make sure you are able to consume from kafka as shown

NiFi-Example-Workflow

5



Step 5 - Extract the contents of flowfile into attributes

Our records contain 3 types of data i.e. factory, warehouse, city based.

e.g.

Sample - 1

```
{  
  "source_type": "factory",  
  "source_name": "Factory A",  
  "temperature": 20.1,  
  "pressure": 1040.88,  
}
```

```
"timestamp": 1727171070
}
```

Sample - 2

```
{
  "source_type": "warehouse",
  "source_name": "Warehouse 1",
  "temperature": 0.37,
  "pressure": 957.9,
  "timestamp": 1727171080
}
```

Sample - 3

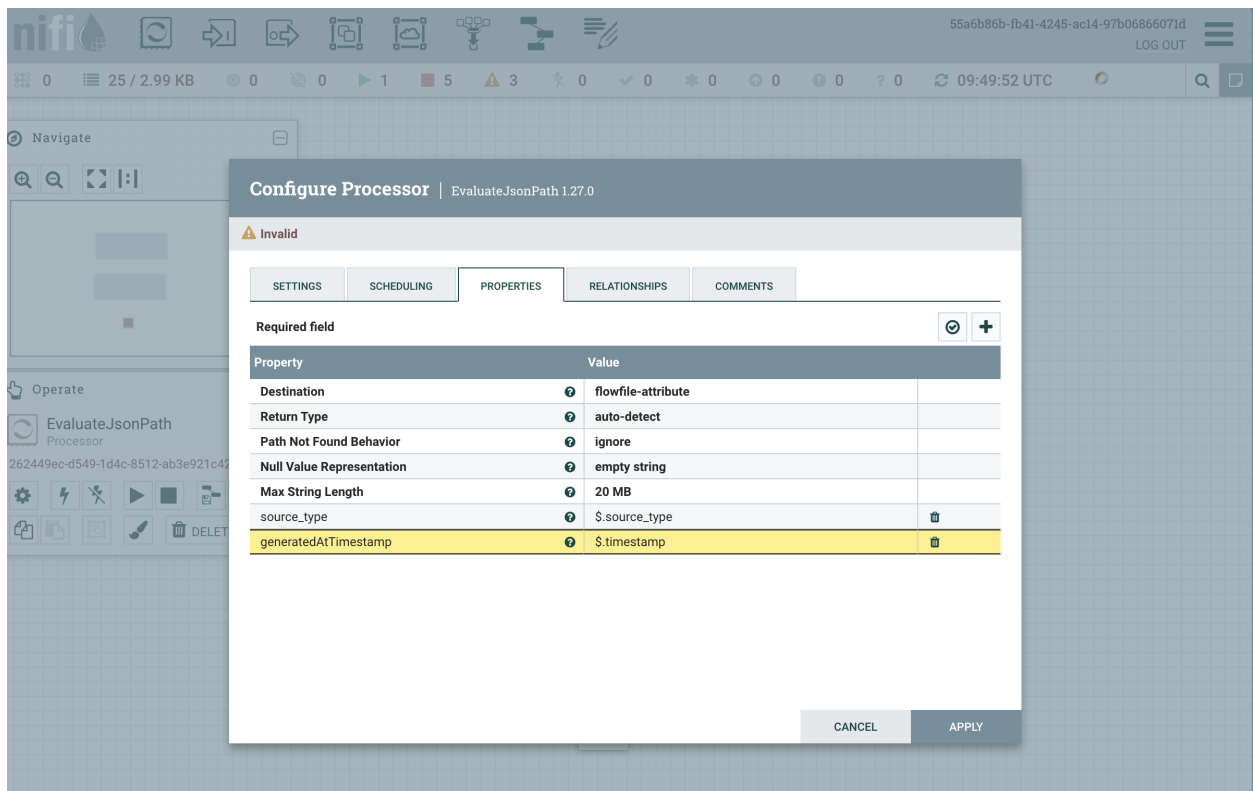
```
{
  "source_type": "city",
  "source_name": "City Center",
  "temperature": 25.23,
  "pressure": 1024.26,
  "timestamp": 1727171105
}
```

Next, we need to extract keys from these json into the flowfile attributes.

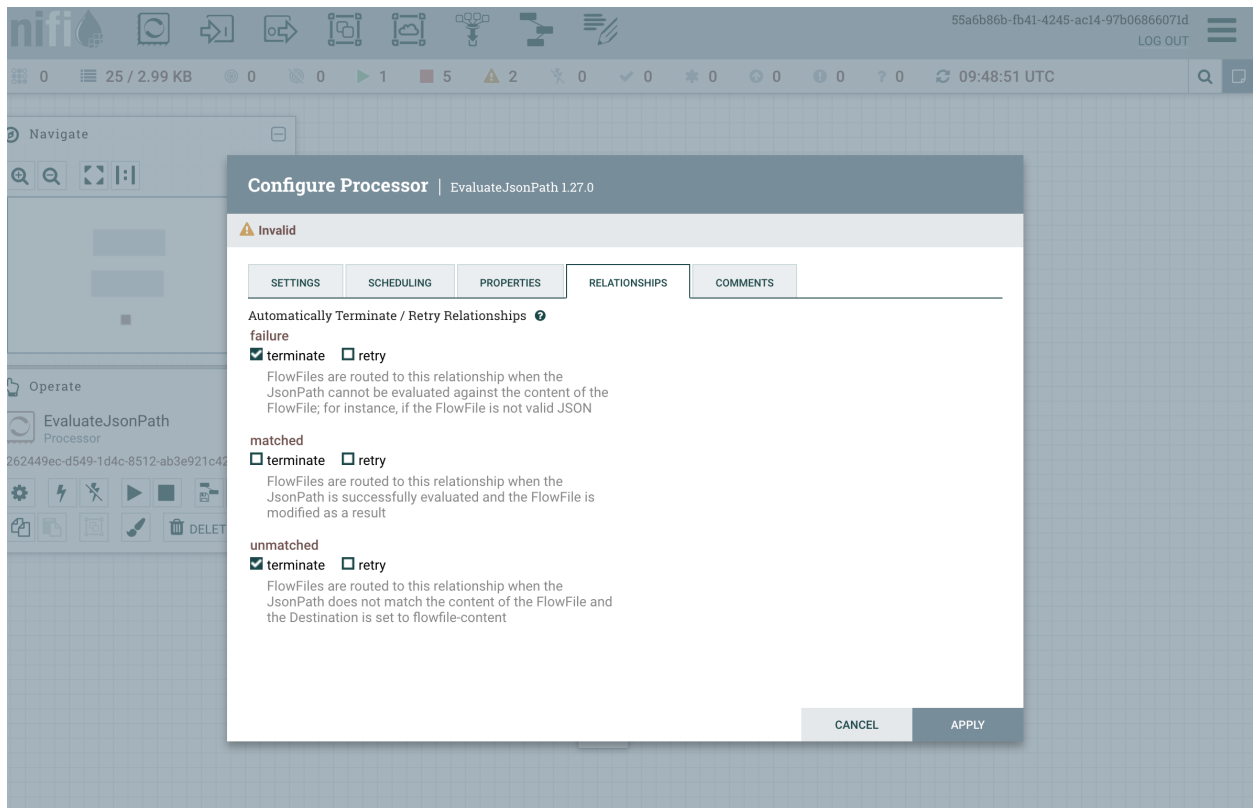
For that, we need to use a processor called as `EvaluateJsonPath`

Let us extract `timestamp` and `source_type` into the flowfile attributes by configuring processor as shown. Add properties manually by clicking + icon in the top right of the processor


Set the destination to `flowfile-attribute`



Terminate the failure and unmatched relationship since we dont need to process those records at all in this example.




Make sure everything looks on UI like shown and there are no errors.

| | |
|---|--|
|  | <div><div>▶</div><div><div>ConsumeKafka_2_6</div><div>ConsumeKafka_2_6 1.27.0</div><div>org.apache.nifi - nifi-kafka-2-6-nar</div></div></div> |
| In | 0 (0 bytes)5 min |
| Read/Write | 0 bytes / 0 bytes5 min |
| Out | 0 (0 bytes)5 min |
| Tasks/Time | 297 / 00:00:03.2005 min |

Name **success**

Queued 0 (0 bytes)

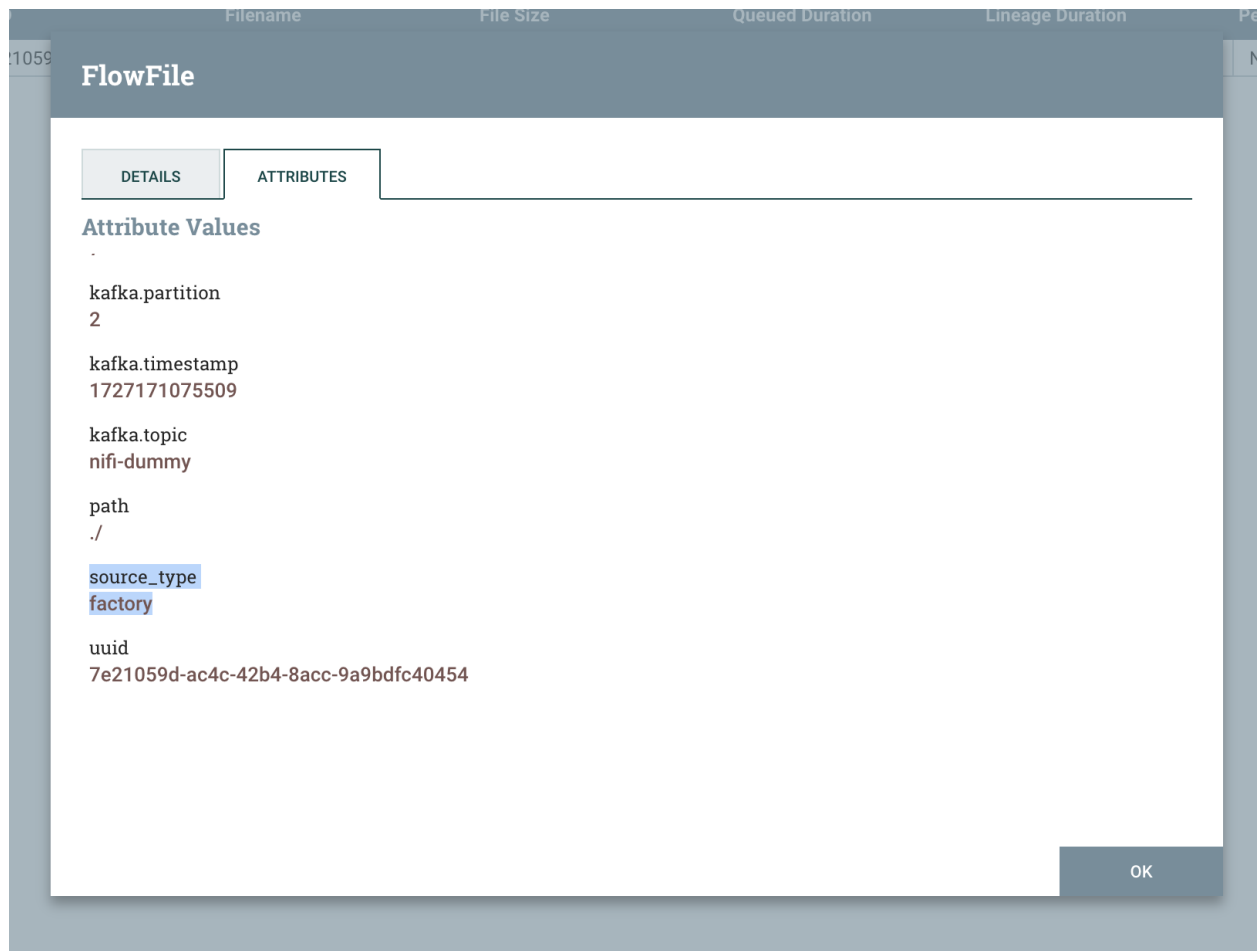
| | | |
|---|---|-------|
|  | EvaluateJsonPath EvaluateJsonPath 1.27.0 org.apache.nifi - nifi-standard-nar | |
| In | 25 (2.99 KB) | 5 min |
| Read/Write | 2.99 KB / 0 bytes | 5 min |
| Out | 25 (2.99 KB) | 5 min |
| Tasks/Time | 25 / 00:00:00.050 | 5 min |

Name **matched**

Queued 25 (2.99 KB)



Also check the attributes of processed flowfiles to see if they contain our custom attributes.



Step 6 - Update the attributes of flowfile to add processedAtTimestamp

Add a processor named UpdateAttributes and configure it as shown. Add a function `${now()}`

Configure Processor | UpdateAttribute 1.27.0

Invalid

SETTINGS

SCHEDULING

PROPERTIES

RELATIONSHIPS

COMMENTS

Required field

✓

+

| Property | | Value | |
|----------------------------------|---|--------------------|--------------|
| Delete Attributes Expression | ? | No value set | |
| Store State | ? | Do not store state | |
| Stateful Variables Initial Value | ? | No value set | |
| Cache Value Lookup Cache Size | ? | 100 | |
| processedAtTimestamp | ? | \${now()} | <div>✕</div> |

⚙️

ADVANCED

CANCEL

APPLY

Make sure it looks good



ConsumeKafka_2_6

ConsumeKafka_2_6 1.27.0

org.apache.nifi - nifi-kafka-2-6-nar

| | | |
|------------|--------------------|-------|
| In | 0 (0 bytes) | 5 min |
| Read/Write | 0 bytes / 0 bytes | 5 min |
| Out | 0 (0 bytes) | 5 min |
| Tasks/Time | 298 / 00:00:03.206 | 5 min |

Name **success**

Queued 0 (0 bytes)



EvaluateJsonPath

EvaluateJsonPath 1.27.0

org.apache.nifi - nifi-standard-nar

| | | |
|------------|-------------------|-------|
| In | 25 (2.99 KB) | 5 min |
| Read/Write | 2.99 KB / 0 bytes | 5 min |
| Out | 25 (2.99 KB) | 5 min |
| Tasks/Time | 25 / 00:00:00.050 | 5 min |

Name **matched**

Queued 0 (0 bytes)



UpdateAttribute

UpdateAttribute 1.27.0

org.apache.nifi - nifi-update-attribute-nar

| | | |
|------------|-------------------|-------|
| In | 25 (2.99 KB) | 5 min |
| Read/Write | 0 bytes / 0 bytes | 5 min |
| Out | 25 (2.99 KB) | 5 min |
| Tasks/Time | 25 / 00:00:00.040 | 5 min |

Name **success**

Queued 25 (2.99 KB)



Check the attributes if it contains our custom attribute.

The screenshot shows the NiFi interface with a FlowFile selected. The 'ATTRIBUTES' tab is active, displaying a list of key-value pairs. The 'processedAtTimestamp' attribute is highlighted, showing the value 'Tue Sep 24 09:54:25 UTC 2024'. Other attributes include 'kafka.timestamp', 'kafka.topic', 'path', 'source_type', 'factory', and 'uuid'. The background shows a table of FlowFiles with columns for UUID, Filename, File Size, Queued Duration, Lineage Duration, and Penalty.

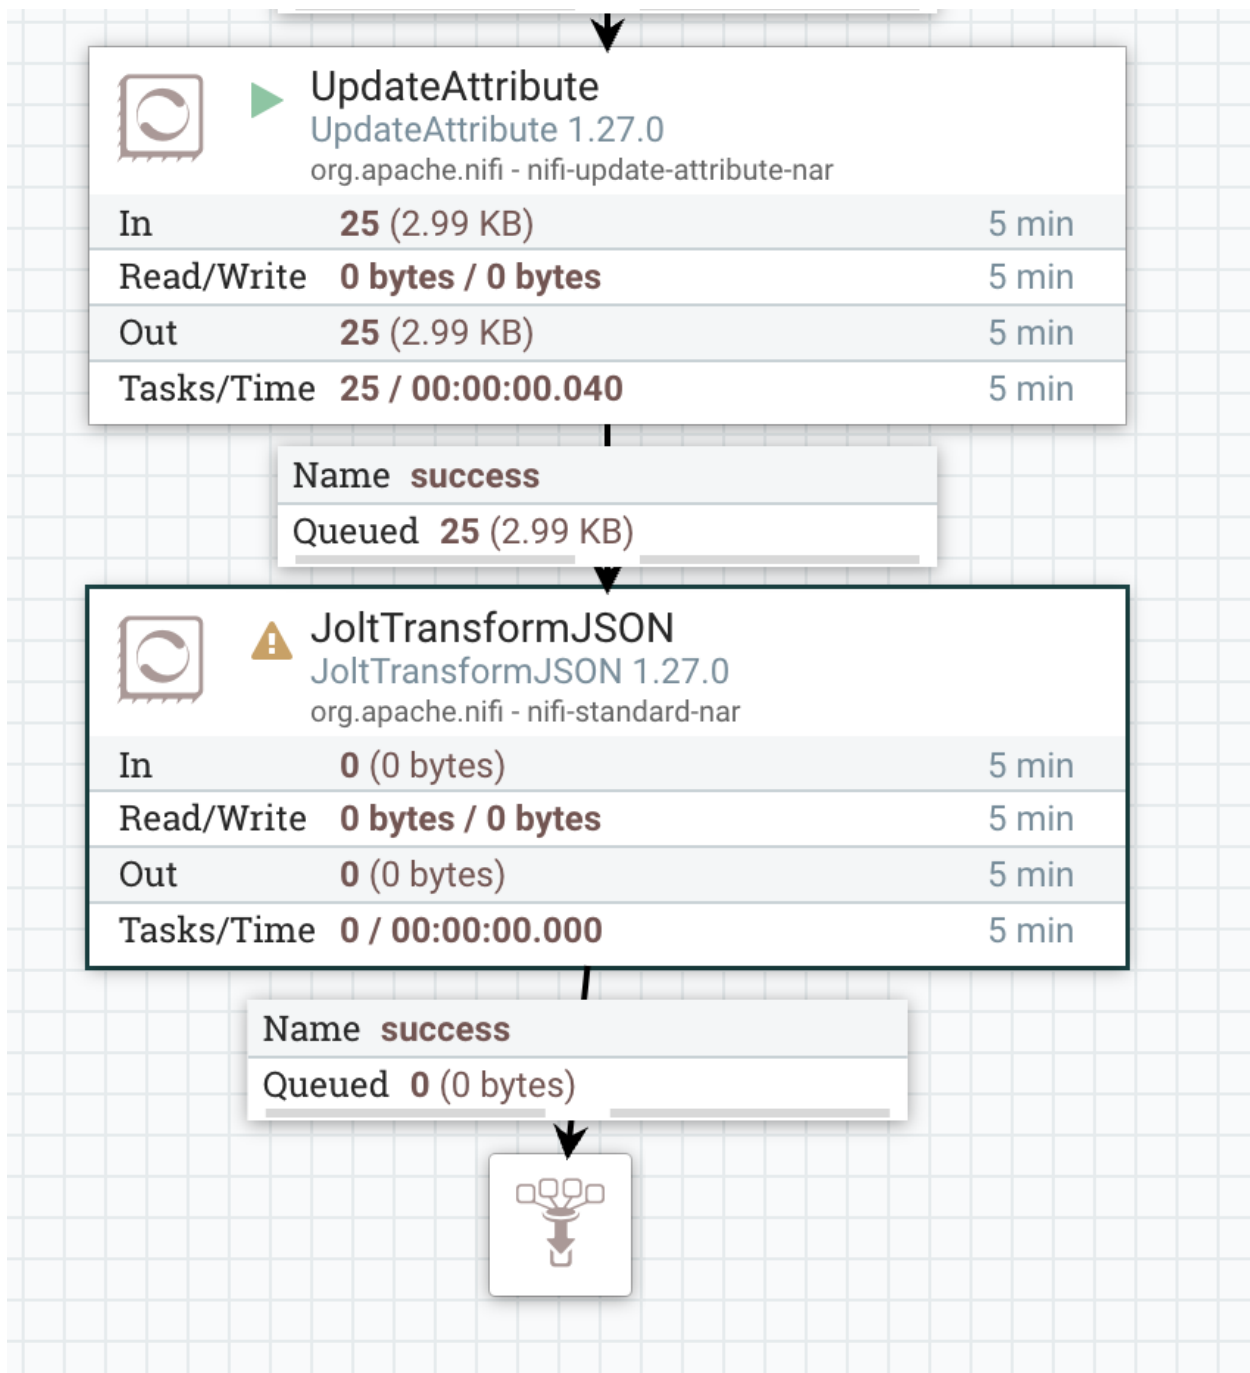
| UUID | Filename | File Size | Queued Duration | Lineage Duration | Penalty |
|-----------------------------|-----------------------------|--------------|-----------------|------------------|---------|
| 7e21059 | | | | | No |
| fe8406a | | | | | No |
| 9f2677e | | | | | No |
| 4a6a8f0 | | | | | No |
| 3a50bde | | | | | No |
| 9b26d93 | | | | | No |
| b6e99d2 | | | | | No |
| 0b68354 | | | | | No |
| d234221 | | | | | No |
| d7b98ec | | | | | No |
| cb691eb | | | | | No |
| 322df4a | | | | | No |
| 9c2fb67 | | | | | No |
| 7c4a514 | | | | | No |
| 18897a4 | | | | | No |
| c540a30 | | | | | No |
| 2b446ac | | | | | No |
| 51bdb4b | | | | | No |
| b4adc9b | | | | | No |
| 20356e3 | | | | | No |
| a1c7558 | | | | | No |
| b519ea5 | | | | | No |
| e40d414 | | | | | No |
| c9dc2194-2f78-40f3-89ba-... | c9dc2194-2f78-40f3-89ba-... | 122.00 bytes | 00:00:25.884 | 00:08:20.681 | No |
| 00401b3e-c760-4b0d-89e4-... | 00401b3e-c760-4b0d-89e4-... | 122.00 bytes | 00:00:25.884 | 00:08:15.600 | No |

The timestamp is in format `Tue Sep 24 09:54:25 UTC 2024` . **We might need our own format sometimes to support different systems.**

In that case, `${now()}` function can be updated as `${now():format('yyyy-MM-dd HH:mm:ss')}`

Step 7 - Rename and pull the values back to flowfile contents using Jolt processor

Use a processor named `JoltTransformRecord` next in line.



Jolt is very powerful processor to manipulate JSON. It required Jolt specification which is an expression language of its own.

To learn more about Jolt, refer <https://jolt-demo.appspot.com/#inception>

In our case, we want to manipulate JSON as shown where

INPUT

```
{
  "source_type": "factory",
  "source_name": "Factory A",
  "temperature": 20.1,
  "pressure": 1040.88,
  "timestamp": 1727171070
}
```

Transformed OUTPUT

```
{
  "SOURCE_TYPE": "factory",
  "SOURCE_NAME": "Factory A",
  "TEMPERATURE": 20.1,
  "PRESSURE": 1040.88,
  "TIMESTAMP": 1727171070
}
```

For development, click the advanced button in the bottom left corner.

Configure Processor | JoltTransformJSON 1.27.0

Invalid

SETTINGS

SCHEDULING

PROPERTIES

RELATIONSHIPS

COMMENTS

Required field

✓

+

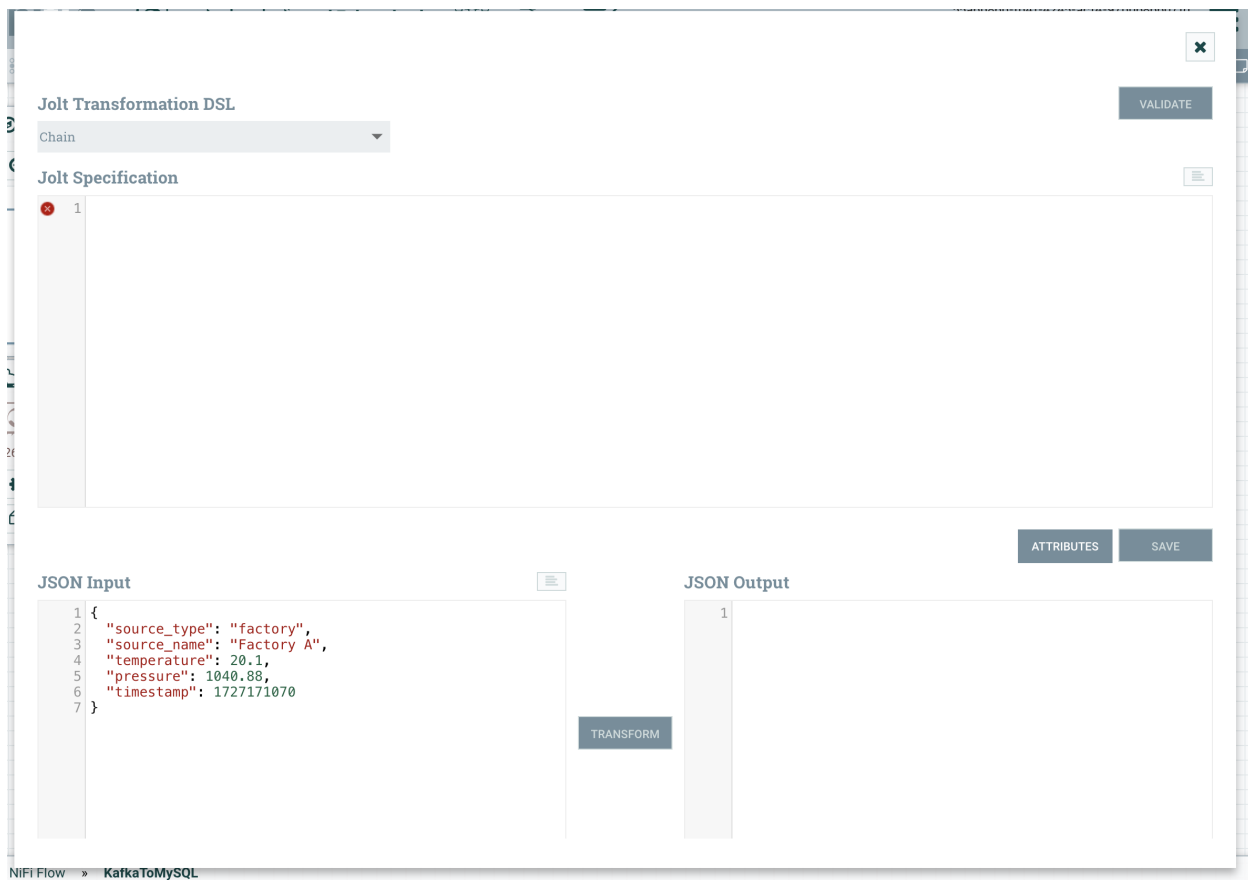
| Property | | Value |
|-------------------------|---|--------------|
| Jolt Transformation DSL | ? | Chain |
| Jolt Specification | ? | No value set |
| Transform Cache Size | ? | 1 |
| Pretty Print | ? | false |
| Max String Length | ? | 20 MB |

⚙️ ADVANCED

CANCEL

APPLY

It will open Jolt DSL window where we can feed our custom JSON and put some specification to test it out.



e.g. to rename our keys, we can put specification

```
[{
  "operation": "shift",
  "spec": {
    "source_type": "SOURCE_TYPE",
    "source_name": "SOURCE_NAME",
    "temperature": "TEMPERATURE",
    "pressure": "PRESSURE",
    "timestamp": "TIMESTAMP"
  }
}]
```

So results will look like,

✕

Jolt Transformation DSL
 Chain

Specification is Valid VALIDATE

Jolt Specification

```

1 [{
2   "operation": "shift",
3   "spec": {
4     "source_type": "SOURCE_TYPE",
5     "source_name": "SOURCE_NAME",
6     "temperature": "TEMPERATURE",
7     "pressure": "PRESSURE",
8     "timestamp": "TIMESTAMP"
9   }
10 }]

```

JSON Input

ATTRIBUTES
SAVE

```

1 {
2   "source_type": "city",
3   "source_name": "City Center",
4   "temperature": 25.23,
5   "pressure": 1024.26,
6   "timestamp": 1727171105
7 }

```

TRANSFORM

JSON Output

```

1 {
2   "SOURCE_TYPE": "city",
3   "SOURCE_NAME": "City Center",
4   "TEMPERATURE": 25.23,
5   "PRESSURE": 1024.26
6 }

```

Now, we have to fetch some extra properties from flowfile attributes. It can be done by adding attribute field to specifications using default operation as shown

```

[ {
  "operation": "shift",
  "spec": {
    "source_type": "SOURCE_TYPE",
    "source_name": "SOURCE_NAME",
    "temperature": "TEMPERATURE",
    "pressure": "PRESSURE",
    "timestamp": "TIMESTAMP"
  }
}, {
  "operation": "default",
  "spec": {
    "EntryProcessedAt": "${processedAtTimestamp}",

```

```

    "KAFKA_TIMESTAMP": "${kafka.timestamp:format('yyyy-MM-dd')}
  }
}]

```

here, `processedAtTimestamp` and `kafka.timestamp` are the names of attribute that we want inside content. We are processing them inside JOLT.

With above specs, our records will look like this

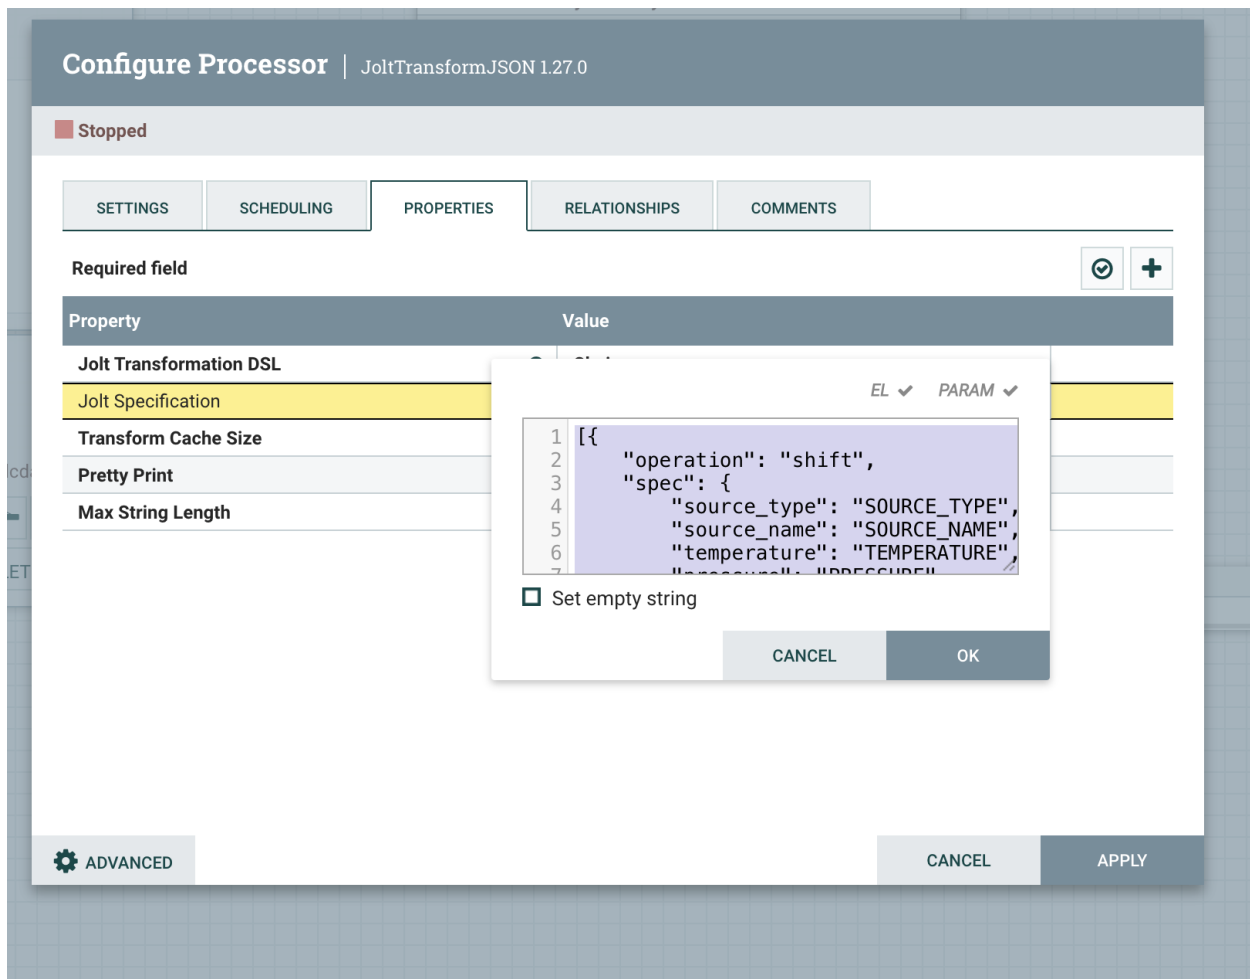
View as: **formatted** ▼

```

1 {
2   "SOURCE_TYPE" : "factory",
3   "SOURCE_NAME" : "Factory A",
4   "TEMPERATURE" : -13.77,
5   "PRESSURE" : 985.06,
6   "TIMESTAMP" : 1727171075,
7   "KAFKA_TIMESTAMP" : "2024-09-24 09:44:35",
8   "EntryProcessedAt" : "Tue Sep 24 09:54:25 UTC 2024"
9 }

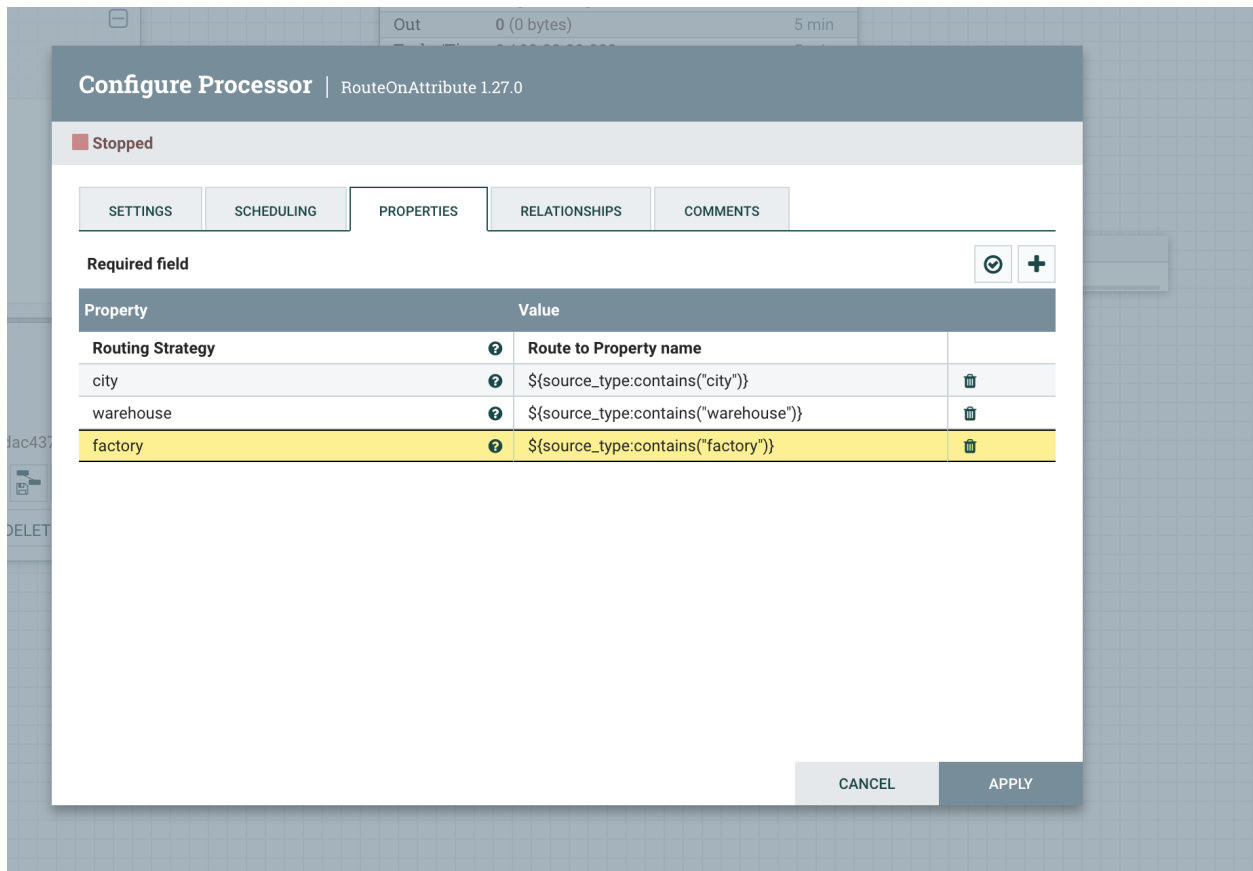
```

Jolt processor, put the spec here.



Step 8 - Route the flowfiles to different destination based on attributes

Add processor named `RouteOnAttribute` and make its properties as shown.



This will essentially create 3 relationships for the processor which will route records accordingly.

Configure Processor | RouteOnAttribute 1.27.0

Invalid

SETTINGS

SCHEDULING

PROPERTIES

RELATIONSHIPS

COMMENTS

Automatically Terminate / Retry Relationships ?

city

☐ terminate ☐ retry

factory

☐ terminate ☐ retry

unmatched

☐ terminate ☐ retry

FlowFiles that do not match any user-define expression will be routed here

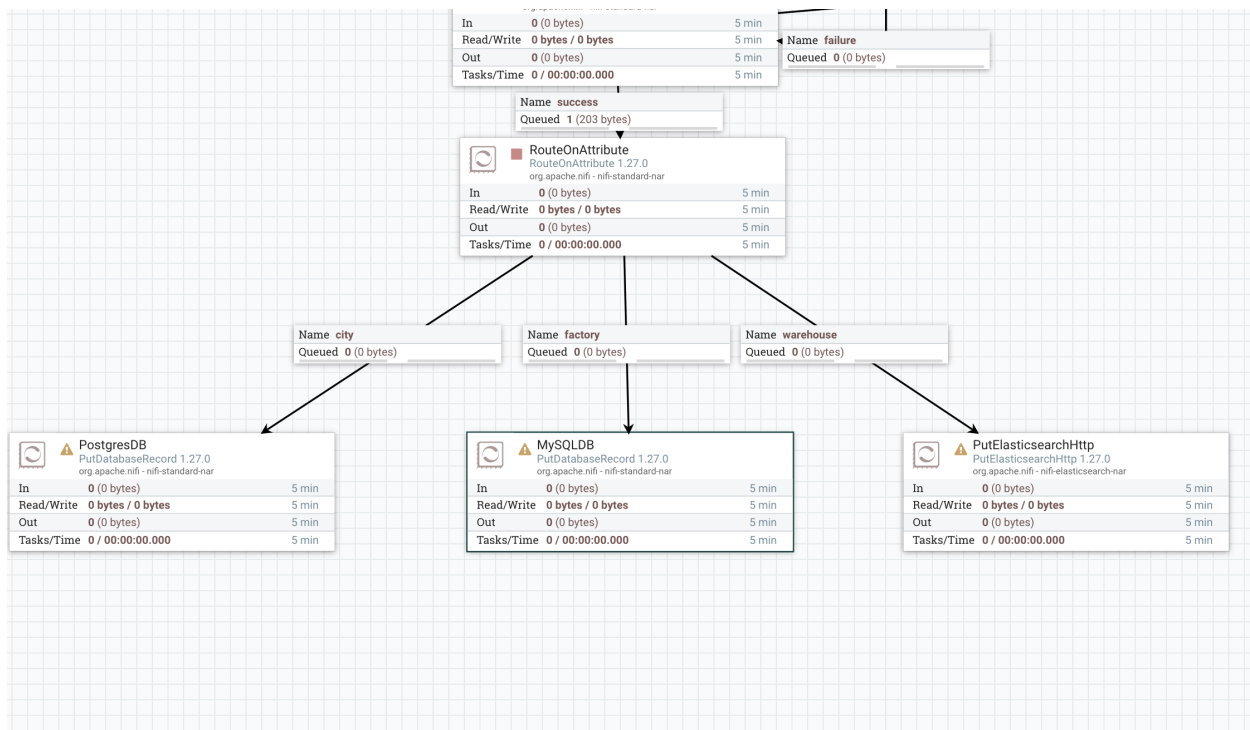
warehouse

☐ terminate ☐ retry

CANCEL

APPLY

Now what we want to achieve here is this -



Step 9 - Configure MySQL and PostgreSQL database connection

Here, we are using `PutDatabaseRecord` processor to send data to MySQL.

Configuration should look like this

Read/Write 0 bytes / 0 bytes5 minName failure
Out 0 (0 bytes)5 minQueued 0 (0 bytes)

Configure Processor

PutDatabaseRecord 1.27.0

Invalid

SETTINGS
SCHEDULING
PROPERTIES
RELATIONSHIPS
COMMENTS

Required field

| Property | Value |
|-------------------------------------|---------------------------|
| Record Reader | JsonTreeReader |
| Database Type | MySQL |
| Statement Type | INSERT |
| Data Record Path | No value set |
| Database Connection Pooling Service | DBCPConnectionPool |
| Catalog Name | No value set |
| Schema Name | No value set |
| Table Name | nifi_dummy |
| Binary String Format | UTF-8 |
| Translate Field Names | true |
| Unmatched Field Behavior | Ignore Unmatched Fields |
| Unmatched Column Behavior | Fail on Unmatched Columns |

CANCEL
APPLY

Here, since we are expecting JSON input, we use JsonTreeReader processor with everything default config.

Then, we have DBCPConnectionPool Controller service which requires extra configuration.

SERVICES

Configure Controller Service | DBCPConnectionPool 1.27.0

INVALID

SETTINGS
PROPERTIES
COMMENTS

Required field

| Property | Value |
|------------------------------|---|
| Database Connection URL | jdbc:mysql://localhost:3306/nifi |
| Database Driver Class Name | com.mysql.jdbc.Driver |
| Database Driver Location(s) | https://repo1.maven.org/maven2/mysql/mys... |
| Kerberos User Service | No value set |
| Kerberos Credentials Service | No value set |
| Kerberos Principal | No value set |
| Kerberos Password | No value set |
| Database User | nifi |
| Password | Sensitive value set |
| Max Wait Time | 500 millis |
| Max Total Connections | 8 |
| Validation query | No value set |

CANCEL
APPLY

We need to put in Database connection URL, Database driver class name, Database driver location jar files, database user and database password

For MySQL

| | |
|-----------------------------|--|
| Database Connection URL | jdbc:mysql://hostname:3306/databasename?serverTimezone=UTC |
| Database Driver Class Name | com.mysql.jdbc.Driver |
| Database Driver Location(s) | https://repo1.maven.org/maven2/mysql/mysql-connector-java/8.0.19/mysql-connector-java-8.0.19.jar |

For PostgreSQL

| | |
|-------------------------|---|
| Database Connection URL | jdbc:postgresql://hostname:3306/databasename?serverTimezone=UTC |
|-------------------------|---|

| | |
|--------------------------------|---|
| Database Driver Class Name | org.postgresql.Driver |
| Database Driver Location(s) | https://jdbc.postgresql.org/download/postgresql-42.7.4.jar |

Now by the design of SQL databases, the tables must be already created before sending data to them. Otherwise, it would complain of table not being existing.

That is it. Try to run entire workflow now and it should send data properly to databases.